



Random Features for Large-Scale Kernel Machines¹

Kailong Wang¹

¹Rutgers University

August 6, 2023

¹Ali Rahimi and Benjamin Recht. “Random Features for Large-Scale Kernel Machines”. In: 20 (2007). Ed. by J. Platt et al. URL: https://proceedings.neurips.cc/paper_files/paper/2007/file/013a006f03dbc5392effeb8f18fda755-Paper.pdf.

Table of Contents

① Motivation


② Random Fourier Features

③ Convergence of RFF



Linear Non-separable Problem


Consider a binary classification problem with non-linear samples.



e.g. For the above dataset $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2]$ where column vector $\mathbf{x}_i \in \mathbb{R}^N$, a linear decision boundary does not exist.

Lifting

One idea is **LIFTING** the samples into a high dimensional space in which the samples are linearly separable.



In this case, the function $\phi(\mathbf{X}) = [\mathbf{x}_1 \circ \mathbf{x}_1, \mathbf{x}_2 \circ \mathbf{x}_2, \sqrt{2}\mathbf{x}_1 \circ \mathbf{x}_2]$, where \circ is the Hadamard product, lifts the samples into \mathbb{R}^3 and the samples are linearly separable.



SVM²

The idea of lifting has been implemented in many classification algorithms such as *support vector machine* (SVM).

Dual Problem of SVM

$$\max_{\alpha} \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N \alpha_n \alpha_m y_n y_m \langle \mathbf{x}_n, \mathbf{x}_m \rangle$$

Dual Problem with Lifting

$$\max_{\alpha} \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N \alpha_n \alpha_m y_n y_m \langle \phi(\mathbf{x}_n), \phi(\mathbf{x}_m) \rangle$$

This is the hard-margin SVM. The soft-margin SVM is similar.

²Stephen Boyd and Lieven Vandenberghe. "Convex optimization". In: (2004).



Curse of Dimensionality–Type I

$$\begin{aligned}\langle \phi(\mathbf{x}_n), \phi(\mathbf{x}_m) \rangle &= [x_{n,1}^2, x_{n,2}^2, \sqrt{2}x_{n,1}x_{n,2}] [x_{m,1}^2, x_{m,2}^2, \sqrt{2}x_{m,1}x_{m,2}]^\top \\ &= x_{n,1}^2 x_{m,1}^2 + x_{n,2}^2 x_{m,2}^2 + 2x_{n,1}x_{n,2}x_{m,1}x_{m,2}\end{aligned}$$

For the given example, it does three multiplication to get the result (a constant value). For a function lifting the original vector space to a much higher dimension, such a calculation can be computationally thirsty. Alternatively, this can be done as follow, whose computational complexity only depends on the dimension of the original vector space.

$$\begin{aligned}(\langle \mathbf{x}_n, \mathbf{x}_m \rangle)^2 &= ([x_{n,1} \ x_{n,2}][x_{m,1} \ x_{m,2}]^\top)^2 \\ &= (x_{n,1}x_{m,1} + x_{n,2}x_{m,2})^2 \\ &= x_{n,1}^2 x_{m,1}^2 + x_{n,2}^2 x_{m,2}^2 + 2x_{n,1}x_{n,2}x_{m,1}x_{m,2} \\ &= \langle \phi(\mathbf{x}_n), \phi(\mathbf{x}_m) \rangle\end{aligned}$$



Kernel Trick

The type of function, such as $(\langle \cdot, \cdot \rangle)^2$, that provides a computationally efficient way to compute the inner product in the high dimensional space is called a **Kernel Function**.

$$K(\cdot, \cdot) = \langle \phi(\cdot), \phi(\cdot) \rangle$$

The matrix that is formed by stacking the kernel function for all samples is called the **Kernel Matrix** or **Gram Matrix K**,

$$\mathbf{K}_{nm} \equiv K(\mathbf{x}_n, \mathbf{x}_m).$$

Some kernel functions can lift the original vector space to an infinite dimensional space. The algorithms involve kernel trick is called **Kernel Machines**.

Curse of Dimensionality–Type II

Another famous kernel machine is kernel ridge regression (KRR). With $\mathbf{y} \in \mathbb{R}^N$, $\mathbf{X} \in \mathbb{R}^{N \times d}$, and $\phi_{d \rightarrow k}(\cdot) : \mathbb{R}^d \rightarrow \mathbb{R}^k$, the loss function is

$$\mathcal{L}(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{y} - \phi(\mathbf{X})\mathbf{w})^\top (\mathbf{y} - \phi(\mathbf{X})\mathbf{w}) + \lambda \mathbf{w}^\top \mathbf{w}.$$

The normal equation of KRR is

$$\begin{aligned} \mathbf{w} &= (\phi(\mathbf{X})^\top \phi(\mathbf{X}) + \lambda \mathbf{I}_k)^{-1} \phi(\mathbf{X})^\top \mathbf{y} \\ &= (\mathbf{K} + \lambda \mathbf{I}_k)^{-1} \phi(\mathbf{X})^\top \mathbf{y}. \end{aligned}$$

Solving this problem requires $\Theta(k^3)$ time and $\Theta(k^2)$ memory.



Motivation

Can we find a way to construct the **Kernel Matrix**, which is equivalent to lift \mathbf{X} to \mathbb{R}^s with $d < s \ll k$, while not sacrifices model performance?

Some Prerequisites

Definition: Shift Invariant Kernel (Radial Basis Function (RBF))

A kernel function $K(\mathbf{x}_n, \mathbf{x}_m)$ is called **shift invariant** if it can be written as $K(\mathbf{x}_n, \mathbf{x}_m) = k(\mathbf{x}_n - \mathbf{x}_m)$ for some function $g(\cdot)$ (e.g. $K_{Gaussian}(\mathbf{x}_n, \mathbf{x}_m) = \exp(-\gamma \|\mathbf{x}_n - \mathbf{x}_m\|_2^2)$).

Mercer's Theorem

A continuous function $K(\mathbf{x}_n, \mathbf{x}_m)$ is a valid kernel function if and only if the kernel matrix \mathbf{K} is **positive semi-definite**.

Bochner's Theorem

A continuous function $k(\cdot)$ is **positive semi-definite** if and only if it is the Fourier transform of a non-negative measure.

Random Fourier Features

Conclusion

A continuous **shift invariant** kernel $K(\mathbf{x}_n, \mathbf{x}_m)$, which is **positive semi-definite** (Mercer's Theorem), is the Fourier transform of a non-negative measure $p(\cdot)$.

$$\phi(\mathbf{x}_n)^\top \phi(\mathbf{x}_m) = K(\mathbf{x}_n, \mathbf{x}_m) = k(\mathbf{x}_n - \mathbf{x}_m) \quad (1)$$

$$= \int_{\mathbb{R}^d} p(\omega) \exp(i\omega^\top (\mathbf{x}_n - \mathbf{x}_m)) d(\mathbf{x}_n - \mathbf{x}_m) \quad (2)$$

$$= \mathbb{E}_\omega [\xi_\omega(\mathbf{x}_n)^\mathsf{H} \xi_\omega(\mathbf{x}_m)] \quad (3)$$

Here $\xi_\omega(\mathbf{x}_n - \mathbf{x}_m) = \exp(i\omega^\top (\mathbf{x}_n - \mathbf{x}_m))$.



Random Fourier Features

Since both the $p(\cdot)$ and $k(\Delta)$ are real-valued, we can replace $\exp(i\omega^\top(\mathbf{x}_n - \mathbf{x}_m))$ with $\cos(\omega^\top(\mathbf{x}_n - \mathbf{x}_m))$. Let

$z_\omega(\mathbf{x}) = \begin{bmatrix} \cos(\mathbf{x}) \\ \sin(\mathbf{x}) \end{bmatrix} = \sqrt{2} \cos(\omega^\top \mathbf{x} + b)$ where ω is drawn from $p(\omega)$ and b is uniformly drawn from $[0, 2\pi]$. Then eq. (3) becomes $\mathbb{E}_\omega[z_\omega(\mathbf{x}_n)^\top z_\omega(\mathbf{x}_m)]$.

To further reduce the variance of the estimator, we can randomly draw s samples of ω and normalize each corresponding $z_\omega(\mathbf{x})$ by \sqrt{s} . Then the inner product $z(\mathbf{x}_n)^\top z(\mathbf{x}_m) = \frac{1}{s} \sum_{j=1}^s z_{\omega_j}(\mathbf{x}_n)^\top z_{\omega_j}(\mathbf{x}_m)$



Algorithm

Algorithm Random Fourier Features

Require: A shift invariant kernel $K(\mathbf{x}_n, \mathbf{x}_m) = k(\mathbf{x}_n - \mathbf{x}_m)$.

Ensure: A randomized feature map $z(\mathbf{x}) : \mathbb{R}^d \rightarrow \mathbb{R}^s$ so that $z(\mathbf{x}_n)^\top z(\mathbf{x}_m) \approx K(\mathbf{x}_n, \mathbf{x}_m)$.

Compute the Fourier transform $p(\cdot)$ of the kernel K : $p(\omega) = \frac{1}{2\pi} \int \exp(-i\omega^\top \Delta) k(\Delta) d\Delta$

Draw s i.i.d. samples $\omega_1, \omega_2, \dots, \omega_s \in \mathbb{R}^d$ from $p(\cdot)$ and s i.i.d. samples $b_1, b_2, \dots, b_s \in [0, 2\pi]$.

Let $z(\mathbf{x}) \equiv \sqrt{\frac{2}{s}} [\cos(\omega_1^\top \mathbf{x} + b_1), \cos(\omega_2^\top \mathbf{x} + b_2), \dots, \cos(\omega_s^\top \mathbf{x} + b_s)]^\top$



Common RFF

Kernel	$K(\Delta)$	$p(\omega)$
Gaussian	$\exp(-\gamma \ \Delta\ _2^2)$	$(2\pi)^{-\frac{s}{2}} \exp(-\gamma \ \omega\ _2^2)$
Laplacian	$\exp(-\ \Delta\ _1)$	$\prod_d (\pi(1 + \omega_d^2))^{-1}$
Cauchy	$\prod_d 2(1 + \Delta_d^2)^{-1}$	$\exp(-\ \omega\ _1)$



Convergence with Hoeffding's Inequality³

Hoeffding's Inequality

Let X_1, X_2, \dots, X_N be independent random variables. Assume that $X_i \in [m_i, M_i]$ for every i . Then, for any $\epsilon > 0$, we have

$$\mathbb{P}\left(\left|\sum_{i=1}^N (X_i - \mathbb{E}[X_i])\right| \geq \epsilon\right) \leq 2 \exp\left(-\frac{2\epsilon^2}{\sum_{i=1}^N (M_i - m_i)^2}\right).$$

Bound for *any* pair of samples \mathbf{x}_n and \mathbf{x}_m

Given z_w is bounded random variable between $[-\sqrt{2/s}, \sqrt{2/s}]$, with Hoeffding's Inequality, we have

$$\mathbb{P}(|z(\mathbf{x}_n)^T z(\mathbf{x}_m) - K(\mathbf{x}_n, \mathbf{x}_m)| \geq \epsilon) \leq 2 \exp\left(-\frac{s\epsilon^2}{4}\right).$$

³Roman Vershynin. "High-Dimensional Probability: An Introduction with Applications in Data Science". In: (2018).

Convergence

Bound for *all* pair of samples \mathbf{x}_n and \mathbf{x}_m

Let \mathcal{M} be a compact subset of \mathbb{R}^d with diameter $\text{diam}(\mathcal{M})$. Then, for the mapping z defined in Algorithm 1, we have

$$\begin{aligned} & \mathbb{P}\left(\sup_{x,y \in \mathcal{M}} |z(\mathbf{x}_n)^\top z(\mathbf{x}_m) - K(\mathbf{x}_n, \mathbf{x}_m)| \geq \epsilon\right) \\ & \leq 2^8 \left(\frac{\sigma_{p(\cdot)} \text{diam}(\mathcal{M})}{\epsilon}\right)^2 \exp\left(-\frac{s\epsilon^2}{4(d+2)}\right). \end{aligned}$$

The $\sigma_{p(\cdot)}^2 = \mathbb{E}_{p(\cdot)}[\omega^\top \omega]$ is the second moment of the Fourier transform of the $K(\cdot, \cdot)$.

The proof of this bound uses the knowledge of ϵ -net and ϵ -covering number.